


# Perceptron






# Some Algebra reminder - Hyperplane


- Hyperplane
  - A subspace whose dimension is one less than that of its ambient space.
  - 1D : A point




- 2D : A line



- 3D : A plane






# Some Algebra reminder

- How do we define a hyperplane in the space?
  - The space of the hyperplane is  $n-1$  (if  $n$  is the space that we work on)
  - All the point on the hyperplane solve the equation
$$\theta_1 x_1 + \dots + \theta_n x_n = b \quad (= \theta_0)$$
    - Where  $x$  are the point coordinates
  - The hyperplane separates the space into two half-spaces
    - All the point that the equation result  $> b$
    - All the point that the equation result  $< b$

# Hyperplane - examples



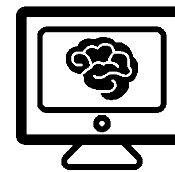
In order to find the distance from the hyperplane we need to normalized  $\theta$ ,  $\theta_0$  by  $\|\theta\|$  (for points on the plane, the normalization doesn't matter, why?)






# Linear separator

- We want to find linear separator:
  - All point above with result greater than 0, will be belong to the +1 class (or -1)
  - All point under with result lower than 0, will be belong to the -1 class (or +1)
- So, what do we need to find?
  - The hyperplane weights  $\theta \in R^{n+1}$  (n hyperplane weights & the bias  $\theta_0$ )
  - We will predict 1 if  $\sum_{i=1}^n \theta_i x_i + \theta_0 > 0$  and -1 otherwise



## Boolean functions – AND

- $X_1$  AND  $X_2$




- Solution?
- If  $1 \times X_1 + 1 \times X_2 - 1.5 > 0$  predict 1
- Otherwise predict -1.
- i.e.  $\theta_0 = -1.5, \theta_1 = 1, \theta_2 = 1$

OR

- $X_1$  OR  $X_2$



- Solution?
- $X_1 + X_2 - 0.5 > 0$  predict 1
- Otherwise -1




# Perceptron



- After we know what we are looking for, linear separator, we need to know how to find it
- Simplest way
  - Start with random weights
  - In each step, improve if necessary – if error exist
  - $\text{Error} = \text{output} - \text{target}$



# Perceptron



- The change in the weight will be

$$\Delta\theta_i = -\eta \sum_{d \in D} (o^{(d)} - t^{(d)}) x_i^{(d)}$$

for each dimension




# Perceptron update rule

$$\Delta\theta_i = -\eta \sum_{d \in D} (o^{(d)} - t^{(d)}) x_i^{(d)}$$

- If  $o^{(d)} - t^{(d)} = 0$ , there is no error = no update

o	t	o-t	$x_i$	$\Delta\theta_i$	$x_i \cdot \theta_i$
-1	+1	<0	>0	>0	increased
-1	+1	<0	<0	<0	increased
+1	-1	>0	>0	<0	decreased
+1	-1	>0	<0	>0	decreased



# Perceptron Algorithm


- The algorithm:
  - Initialize weights to some small random number
  - Repeat until convergence (no error = no weight update):
    - For each  $x^{(d)}$  in  $D$  compute: ( $* x^{(d)} = \bar{x}_d$ ):
      - $o^{(d)} = \text{sgn}(\theta \cdot x^{(d)})$
    - For each  $\theta_i$  do:
      - $\Delta\theta_i = -\eta \sum_{d \in D} (o^{(d)} - t^{(d)}) x_i^{(d)}$  for each  $i$
      - Update  $\theta_i = \theta_i + \Delta\theta_i$

# Stochastic Perceptron




- The algorithm:
  - Set weights randomly
  - Repeat until convergence:
    - Choose  $d$  randomly (or in some order)
    - Calculate  $o^{(d)} = \text{sgn}(\theta \cdot x^{(d)})$
    - Calculate  $\Delta\theta_i = -\eta(o^{(d)} - t^{(d)})x_i^{(d)}$  for each  $i$
    - Then update  $\theta_i = \theta_i + \Delta\theta_i$

# Perceptron Example - Stochastic




- Training data:
  - $(-1,-1) \rightarrow +1$ ,  $(-1,+1) \rightarrow +1$ ,  $(+1,-1) \rightarrow +1$ ,  $(+1,+1) \rightarrow -1$
- Weight init:
  - $\theta_0 = 0.1$ ,  $\theta_1 = -0.2$ ,  $\theta_2 = 0.15$ ,  $\eta = 0.05$

# Perceptron Example - Stochastic




- Check  $(-1, -1) \rightarrow +1$
- $\text{sgn}(\theta \cdot x^{(d)}) = 0.1 - 0.2 * (-1) + 0.15 * (-1) = 0.15 > 0$
- $o = +1$
- Since  $t=o$  no update required

# Perceptron Example - Stochastic



- Check  $(-1, +1) \rightarrow +1$
- $\text{sgn}(\theta \cdot x^{(d)}) = 0.1 - 0.2 * (-1) + 0.15 * (+1) = 0.45 > 0$
- $o = +1$
- Since  $t=o$  no update required


# Perceptron Example - Stochastic



- Check  $(+1, -1) \rightarrow +1$
- $\text{sgn}(\theta \cdot x^{(d)}) = 0.1 - 0.2 * (+1) + 0.15 * (-1) = -0.25 < 0$
- $o = -1$
- Since  $t \neq o$  update required:
  - $\theta_{0(\text{new})} = 0.1 - 0.05 * (-1 - 1) * 1 = 0.2$
  - $\theta_{1(\text{new})} = -0.2 - 0.05 * (-1 - 1) * 1 = -0.1$
  - $\theta_{2(\text{new})} = 0.15 - 0.05 * (-1 - 1) * (-1) = 0.05$




# Perceptron Example - Stochastic




- Check  $(+1, +1) \rightarrow -1$
- $\text{sgn}(\theta \cdot x^{(d)}) = 0.2 - 0.1 * (+1) + 0.05 * (+1) = 0.15 > 0$
- $o = +1$
- Since  $t \neq o$  update required:
  - $\theta_{0(\text{new})} = 0.2 - 0.05 * (+1 - (-1)) * 1 = 0.1$
  - $\theta_{1(\text{new})} = -0.1 - 0.05 * (+1 - (-1)) * 1 = -0.2$
  - $\theta_{2(\text{new})} = 0.05 - 0.05 * (+1 - (-1)) * 1 = -0.05$

# Perceptron Example - Stochastic



- Check  $(+1, -1) \rightarrow +1$
- $\text{sgn}(\theta \cdot x^{(d)}) = 0.1 - 0.2 * (+1) - 0.05 * (-1) = -0.05 < 0$
- $o = -1$
- Since  $t \neq o$  update required:
  - $\theta_{0(\text{new})} = 0.1 - 0.05 * (-1 - 1) * 1 = 0.2$
  - $\theta_{1(\text{new})} = -0.2 - 0.05 * (-1 - 1) * 1 = -0.1$
  - $\theta_{2(\text{new})} = -0.05 - 0.05 * (-1 - 1) * (-1) = -0.15$

# Perceptron Example - Stochastic



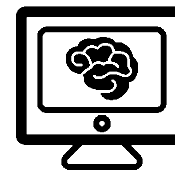
- We got the linear separator:

$$\bar{\theta} \cdot \bar{x} = -0.1 * x_1 - 0.15 * x_2 + 0.2 = 0$$

# Perceptron problem



- What is the problem in the perceptron algorithm?
- How can we solve it?



# LMS – Least Mean Squares

- The algorithm:
  - Initialize weights to some small random number
  - Repeat until convergence (no error = no weight update):
    - For each  $x^{(d)}$  in D compute: ( $* x^{(d)} = \bar{x}_d$ ):
      - $o^{(d)} = (\theta \cdot x^{(d)})$
    - For each  $\theta_i$  do:
      - $\Delta\theta_i = -\eta \sum_{d \in D} (o^{(d)} - t^{(d)}) x_i^{(d)}$
      - Update  $\theta_i = \theta_i + \Delta\theta_i$

# LMS – Least Mean Squares



- Our goal here is not to find linear hyperplane that separate the data, but to minimize the distance from the offset (either +1 or -1) of the hyperplane
- The linear hyperplane that we will use to predict new instance is the outcome of this minimization

# LMS – Least Mean Squares



$$E[\vec{\theta}] = \frac{1}{2} \sum_{d \in D} (o^{(d)} - t^{(d)})^2 = \frac{1}{2} \left[ \sum_{d \in D^+} (o^{(d)} - 1)^2 + \sum_{d \in D^-} (o^{(d)} + 1)^2 \right]$$



*Minimize the distance  
between the positive  
instances and the +1  
iso-line of the function*




*Minimize the distance  
between the negative  
instances and the -1  
iso-line of the function*

# LMS – Least Mean Squares



$$E[\vec{\theta}] = \frac{1}{2} \sum_{d \in D} (o^{(d)} - t^{(d)})^2 = \frac{1}{2} \left[ \sum_{d \in D^+} (o^{(d)} - 1)^2 + \sum_{d \in D^-} (o^{(d)} + 1)^2 \right]$$





# LMS – Least Mean Squares



- What is the differences between Perceptron and LMS?
  - The target is numerical value (not class) – in this case, +1 or -1
  - The output calculation – the result is a number & not a class (the real value of  $\theta \cdot x^{(d)}$  and not  $\text{sgn}(\theta \cdot x^{(d)})$ )
  - The optimization function – in LMS minimum distance from +1 & -1 hyper planes. In perceptron zero classification error
  - LMS will converge, perceptron only if the data are linear separable (no error on training data)



# Logistic Regression

- Linear Regression - Recap
- Predict a continuous value

- Hypothesis function :

$$h_{\theta}(x) = \theta^T x$$

- Cost function (MSE):

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

- Goal :

$$\min_{\theta} J(\theta)$$

$x_1$	$y$
	10000
	21011
	19213
	15213
	18212
	22001



# Linear Regression -> Classification

- What if instead of predicting a continuous value we try to predict a class?

- Hypothesis function :

$$h_{\theta}(x) = \theta^T x$$

- Cost function (MSE):

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

- Goal :

$$\min_{\theta} J(\theta)$$

$x_1$	$y$
	1
	1
	0
	1
	0
	1



# Linear Regression -> Classification

- What if instead of predicting a continuous value we try to predict a class?

- Hypothesis function :

$$h_{\theta}(x) = \theta^T x$$

- Predict


*if  $h_{\theta}(x) > 0.5$  then 1  
else 0*

$x_1$	$y$
	1
	1
	0
	1
	0
	1


# Usually a Bad Idea




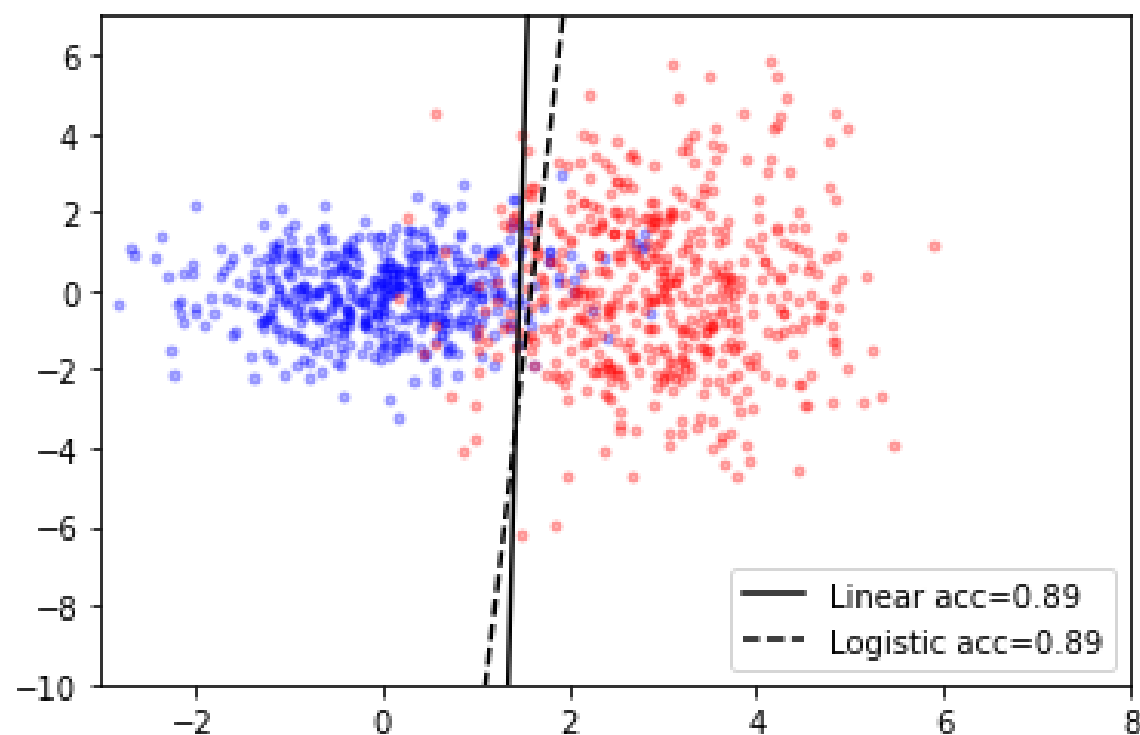

Works ok



Doesn't Work



# Usually a Bad Idea





# Logistic Regression


- We will try to predict the probability the instance belongs to class 1.

- Hypothesis function :

$$h_{\theta}(x) = P(1|x)$$

- How do we go from score to probability?

$x_1$	$y$
	1
	1
	0
	1
	0
	1



# Score to Probability

- Sigmoid function for logistic regression:

$$S(\theta^T x) = \frac{1}{1 + e^{-\theta^T x}} = \frac{e^{\theta^T x}}{1 + e^{\theta^T x}}$$






# Logistic Regression Function

- Sigmoid Function:

$$S(x) = \frac{1}{1 + e^{-x}} = \frac{e^x}{1 + e^x}$$


Score	$-\infty$	-2	0	+2	$+\infty$
Sigmoid (Score)	$\frac{1}{1 + e^{\infty}}$ = 0	$\frac{1}{1 + e^2}$ = 0.12	$\frac{1}{1 + e^0}$ = 0.5	$\frac{1}{1 + e^{-2}}$ = 0.88	$\frac{1}{1 + e^{-\infty}}$ = 1
	?	?	?	?	?



# Logistic Regression Function

- Sigmoid Function:

$$S(x) = \frac{1}{1 + e^{-x}} = \frac{e^x}{1 + e^x}$$





# Logistic Regression Classification

- What if instead of predicting a continuous value we try to predict a class?

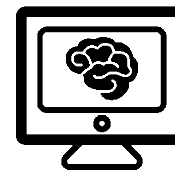
- Hypothesis function :

$$h_{\theta}(x) = S(\theta^T x)$$

- Prediction :

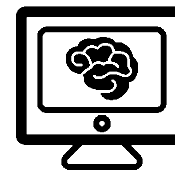
*if  $h_{\theta}(x) > 0.5$  then 1  
else 0*

$x_1$	$y$
	1
	1
	0
	1
	0
	1



# Logistic Regression - Loss

- We have data and we defined the hypothesis function  $h_{\theta}(x)$ .
- We need to find a cost function so we can improve  $h_{\theta}(x)$ .
- We will use Maximum likelihood to find the appropriate cost function.



# Logistic Regression

- $P(y|x, \theta) = (h_{\theta}(x))^y \cdot (1 - h_{\theta}(x))^{1-y}$

- Remember :

$h_{\theta}(x)$  = probability  $x$  belongs to class 1

- And so we get :

- $P(0|x, \theta) = 1 - h_{\theta}(x)$
- $P(1|x, \theta) = h_{\theta}(x)$

# Logistic Regression - ML



$$P(y|x, \theta) = (h_{\theta}(x))^y \cdot (1 - h_{\theta}(x))^{1-y}$$

Assuming independent instances

$$P(D|\theta) \overset{\uparrow}{=} \prod_{d=1}^m P(y^{(d)} | x^{(d)}, \theta) =$$

$$\prod_{d=1}^m (h_{\theta}(x^{(d)}))^y \cdot (1 - h_{\theta}(x^{(d)}))^{1-y}$$

# Logistic Regression - ML



$$\operatorname{argmax}_{\theta} \prod_{d=1}^m \left( h_{\theta}(x^{(d)}) \right)^{y^{(d)}} \cdot \left( 1 - h_{\theta}(x^{(d)}) \right)^{1-y^{(d)}} =$$

$$\operatorname{argmax}_{\theta} \ln \left( \prod_{d=1}^m \left( h_{\theta}(x^{(d)}) \right)^{y^{(d)}} \cdot \left( 1 - h_{\theta}(x^{(d)}) \right)^{1-y^{(d)}} \right) =$$

$$\operatorname{argmax}_{\theta} \sum_{d=1}^m \ln \left( \left( h_{\theta}(x^{(d)}) \right)^{y^{(d)}} \right) + \ln \left( \left( 1 - h_{\theta}(x^{(d)}) \right)^{1-y^{(d)}} \right) =$$

$$\operatorname{argmax}_{\theta} \sum_{d=1}^m y^{(d)} \cdot \ln \left( h_{\theta}(x^{(d)}) \right) + (1 - y^{(d)}) \cdot \ln \left( 1 - h_{\theta}(x^{(d)}) \right) =$$

$$\operatorname{argmin}_{\theta} \sum_{d=1}^m -y^{(d)} \cdot \ln \left( h_{\theta}(x^{(d)}) \right) - (1 - y^{(d)}) \cdot \ln \left( 1 - h_{\theta}(x^{(d)}) \right)$$


# Logistic Regression Cost Function



## Binary Cross Entropy



$$\frac{1}{m} \sum_{d=1}^m -y^{(d)} \cdot \ln \left( h_{\theta}(x^{(d)}) \right) - (1 - y^{(d)}) \cdot \ln \left( 1 - h_{\theta}(x^{(d)}) \right)$$





# Cost Function Intuition

- $Cost(x, \theta) = \begin{cases} -\log(h_{\theta}(x)) & y = 1 \\ -\log(1 - h_{\theta}(x)) & y = 0 \end{cases}$



# Learn Logistic Regression



$$\text{cost}(\vec{\theta}) = - \sum_{d=1}^m y^{(d)} \ln \left( s(\theta, \vec{x}^{(d)}) \right) + (1 - y^{(d)}) \ln \left( 1 - s(\vec{\theta}, \vec{x}^{(d)}) \right)$$

$$\ln(s(\vec{\theta}, \vec{x}^{(d)})) = \ln \left( \frac{1}{1 + e^{-\theta^T x^{(d)}}} \right) = -\ln \left( 1 + e^{-\theta^T x^{(d)}} \right)$$

$$\begin{aligned} \ln(1 - s(\vec{\theta}, \vec{x}^{(d)})) &= \ln \left( 1 - \frac{1}{1 + e^{-\theta^T x^{(d)}}} \right) \\ &= \ln \left( \frac{1 + e^{-\theta^T x^{(d)}}}{1 + e^{-\theta^T x^{(d)}}} - \frac{1}{1 + e^{-\theta^T x^{(d)}}} \right) = \ln \left( \frac{e^{-\theta^T x^{(d)}}}{1 + e^{-\theta^T x^{(d)}}} \right) \\ &= -\theta^T x^{(d)} - \ln \left( 1 + e^{-\theta^T x^{(d)}} \right) \end{aligned}$$

# Learn Logistic Regression

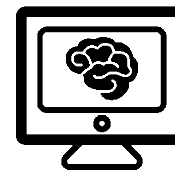


$$\text{cost}(\vec{\theta}) = - \sum_{d=1}^m -y^{(d)} \ln \left( 1 + e^{-\theta^T x^{(d)}} \right) + (1 - y^{(d)}) \left( -\theta^T x^{(d)} - \ln \left( 1 + e^{-\theta^T x^{(d)}} \right) \right)$$

$$= - \sum_{d=1}^m y^{(d)} \theta^T x^{(d)} - \theta^T x^{(d)} - \ln \left( 1 + e^{-\theta^T x^{(d)}} \right) =$$

$$- \sum_{d=1}^m y^{(d)} \theta^T x^{(d)} - \theta^T x^{(d)} - \ln \left( \frac{1 + e^{\theta^T x^{(d)}}}{e^{\theta^T x^{(d)}}} \right)$$

$$- \sum_{d=1}^m y^{(d)} \theta^T x^{(d)} - \ln \left( 1 + e^{\theta^T x^{(d)}} \right)$$



# Learn Logistic Regression

$$\text{cost}(\vec{\theta}) = - \sum_{d=1}^m y^{(d)} \theta^T x^{(d)} - \ln \left( 1 + e^{\theta^T x^{(d)}} \right)$$

- The derivation for each instance will give:

$$\frac{\partial}{\partial \theta_i} \text{cost}(\vec{x}, \vec{\theta}) = - \left( y - S(\vec{\theta}, \vec{x}) \right) x_i$$

- i – is the i feature
- And for all m training data:

$$\frac{\partial}{\partial \theta_i} \text{cost}(\vec{\theta}) = \sum_{d=1}^m \left( S(\vec{\theta}, \vec{x}^{(d)}) - y^{(d)} \right) x_i^{(d)}$$

- You can now use gradient descent for finding the best  $\vec{\theta}$

# Logistic Regression Summary



- Hypothesis function :

$$h_{\theta}(x) = S(\theta^T x)$$

- Cost function:

$$J(\theta) = \frac{1}{m} \sum_{d=1}^m -y^{(d)} \cdot \ln(h_{\theta}(x^{(d)})) - (1 - y^{(d)}) \cdot \ln(1 - h_{\theta}(x^{(d)}))$$

- Goal :

$$\min_{\theta} J(\theta)$$

# Multi-class classification



- How can we convert 2 classes linear separator to solve multi-class problem?
- We will create a predictor for each class (one vs. all)
- Predict  $i$  if  $f_i(x) > f_j(x) \forall j \neq i$

# Questions

