




Dimensionality Reduction





Agenda

- Why?
- Methods
 - Feature selection
 - Filter
 - Wrapper
 - Embedded
 - Feature extraction
 - PCA




Why Reduce Dimensionality?

- The computation perspective
 - Reduces time complexity – Less computation
 - Reduces space complexity – Less parameters
- The modeling perspective
 - Information is less scattered – in a high dimension everything is far
 - Reduces sample complexity
 - More interpretable; simpler explanation

Why Reduce Dimensionality?



- The curse of dimensionality cause the data to be sparse in high features space
- In order to maintain the density we need to increase the training example exponentially
- In practice, the curse of dimensionality means that, for a given sample size, there is a maximum number of features above which the performance of our classifier will degrade rather than improve



Example I




X1	X2	X3	X4	Y
10	14	24	1	1
50	2	5	5	1
5	16	30	0.5	2
30	10	2	3	3
10	23	21	1	2



Example II

- Consider a 3-class classification problem
- Simple approach:
- Divide the feature space into uniform bins
- Compute the ratio of examples for each class at each bin
- For a new example, find its bin and choose the predominant class in that bin
- In 1 feature dimension with 9 samples:

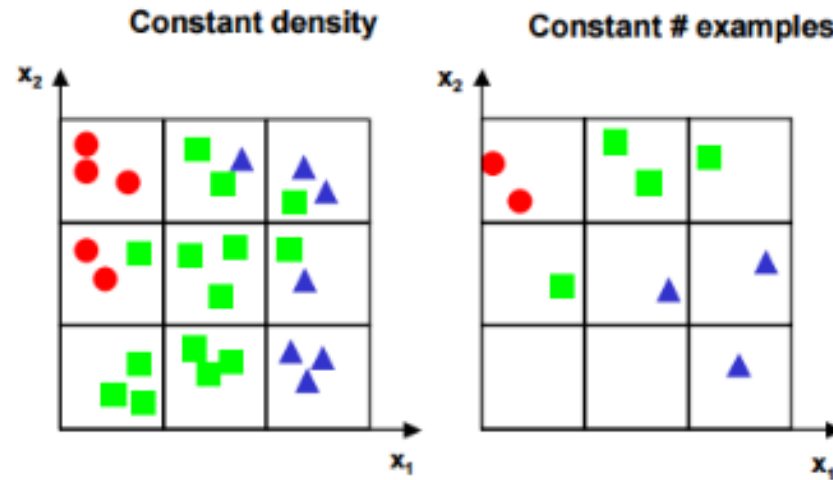


- But there is too much overlap – bad classification

Example II



- We decide to incorporate a second feature (9 bins instead of 3):




- Maintain the number of examples results in a 2D scatter plot that is very sparse
- Maintain the density increases the number of examples from 9 (in 1D) to 27 (in 2D)

Example II



- Moving to three features makes the problem worse:
 - The number of bins grows to $3^3=27$
 - To achieve the same density we need 81 examples
 - For the same number of examples, well, the 3D scatter plot is almost empty



Reducing features space



- We will talk about 2 processes:


- Feature selection

- The process of selecting a subset of relevant features for use in model construction



- Feature extraction

- Creates new set of features that better represent the data (usually in lower dimension)



Feature selection



- We want to find a way to select features in reasonable time (polynomial)
- 3 approaches for feature selection:
 - Wrapper
 - In every step try to learn with a subset of features and select the subset with the best accuracy (less error)
 - Must be greedy selection – otherwise will be exponentially
 - Filter
 - Run a test for each feature and select the ones with the highest scores
 - Embedded
 - During the learning phase the algorithm calculate 'importance score'



Feature selection - Wrapper

- How many feature subset possible? 2^n
- Forward
 - Start with empty set
 - In each iteration add the feature that most increases the accuracy
 - Stop when you reach a predefined accuracy / number of features
- Backward
 - Start with a full set
 - In each iteration remove the feature that minimally reduces the accuracy
 - Stop when you reach a predefined accuracy / number of features

Feature selection - Wrapper



- More search strategies
 - Greedy search on part of the space, exponential search on the other (small part)
 - Hybrid – Plus L Minus R
 - Bidirectional
 - Random

Feature selection - Filter




- Select subset of variables, independently of the learning model
- Provide generic selection of features, not tuned by a given learner (universal) – it can be drawback, selected features are not optimized for the used classifier

Feature selection - Filter



- One of the most basic idea is to find correlation to the target function
 - Pearson Correlation – Measures linear dependency
 - Mutual Information – Measures non-linear dependency



Feature selection - Filter

- Pearson Correlation

$$\rho(x_k, y) = \frac{\sum_{d=1}^m (x_k^{(d)} - \mu_k)(y^{(d)} - \mu_y)}{\sqrt{\sum_{d=1}^m (x_k^{(d)} - \mu_k)^2 \sum_{d=1}^m (y^{(d)} - \mu_y)^2}} = \frac{\sigma_{x_k y}}{\sigma_{x_k} \sigma_y}$$

- $x_k^{(d)}$ – the d value of the feature k
- $y^{(d)}$ – the d value of the function target y
- μ_k – the mean of the feature k
- μ_y – the mean of the function target y



Pearson Correlation – example I

$$\rho = \frac{\sum_{d=1}^m (x_k^{(d)} - \mu_k)(y^{(d)} - \mu_y)}{\sqrt{\sum_{d=1}^m (x_k^{(d)} - \mu_k)^2 \sum_{d=1}^m (y^{(d)} - \mu_y)^2}} = \frac{\sigma_{x_k y}}{\sigma_{x_k} \sigma_y}$$

X_k	Y
0.5377	0
1.8339	0
-2.2588	1
0.8622	1
0.3188	0
-1.3077	0
-0.4336	0
0.3426	1
3.5784	1
2.769	1

$$\begin{aligned}\mu_k &= 0.6243 \\ \mu_y &= 0.5\end{aligned}$$

$X_k - \mu_k$	$Y - \mu_y$
-0.0866	-0.5
1.2096	-0.5
-2.8831	0.5
0.2379	0.5
-0.3055	-0.5
-1.932	-0.5
-1.0579	-0.5
-0.2817	0.5
2.9541	0.5
2.1447	0.5

$$\rho = 0.2587$$



Pearson Correlation – example II

$$\rho = \frac{\sum_{d=1}^m (x_k^{(d)} - \mu_k)(y^{(d)} - \mu_y)}{\sqrt{\sum_{d=1}^m (x_k^{(d)} - \mu_k)^2 \sum_{d=1}^m (y^{(d)} - \mu_y)^2}} = \frac{\sigma_{x_k y}}{\sigma_{x_k} \sigma_y}$$

X	Y
0.1	0
0.2	0
0.3	0
0.4	0
0.5	0
0.6	1
0.7	1
0.8	1
0.9	1
1	1

X	Y
1	0
2	0
3	0
4	0
5	0
6	1
7	1
8	1
9	1
10	1

X	Y
-1	0
-2	0
-3	0
-4	0
-5	0
-6	1
-7	1
-8	1
-9	1
-10	1

X	Y
-1	1
-2	1
-3	1
-4	1
-5	1
-6	0
-7	0
-8	0
-9	0
-10	0

$$\rho_1 = 0.8704$$

$$\rho_2 = 0.8704$$

$$\rho_3 = -0.8704$$

$$\rho_4 = 0.8704$$

Feature selection - Filter







- Pearson Correlation
 - 1 – total positive correlation
 - 0 – no correlation
 - -1 – is total negative correlation


Pearson Correlation



Anscombe Quadrant -- Correlation Demonstration



X



From Pearson Correlation to Spearman rank correlation



- Pearson correlation is very sensitive to the actual values of the data and to outliers
- Pearson (and its relatives) measure linearity of the data. This is not always the correct model or desired observation/assessment.
(a strong non-linear relationship can exist but Pearson might not get too excited...)
- It is often much more robust to use rank based correlation measures



Spearman rank correlation

- Transform the data into ranks
- Calculate the correlation on the ranks
- What are the possible values? $-1 \leq SRC \leq 1$


X_k	Y
0.5377	0
1.8339	0
-2.2588	1
0.8622	1
0.3188	0
-1.3077	0
-0.4336	0
0.3426	1
3.5784	1
2.769	1



X_k	Y
6	0
8	0
1	1
7	1
4	0
2	0
3	0
5	1
10	1
9	1

$$\rho = 0.2587$$

$$r_s = 0.3133$$



Embbeded



- Which algorithm can calculate feature importance during the learning?
 - Decision Trees – weighted goodness of split
 - Regression – when the features have the same scale

Feature extraction



- We want to reduce the dimension, but to preserve most of the information
- In other word, we want to take the original features space, probably with correlation between the features, and summarizes it by uncorrelated axes
- First, how we measure information?
 - One way is the Variance
- So, we want to project the data to new axes and preserve the variance

Feature extraction - projection



- A projection is a transformation of data points from one axes system to another
 - Translate – Subtracting the mean vector from all data points, and this is equivalent to moving the center of the data to the origin

$$\begin{pmatrix} x_{1,1} - \mu_1 & \cdots & x_{1,d} - \mu_d \\ \vdots & \ddots & \vdots \\ x_{N,1} - \mu_1 & \cdots & x_{N,d} - \mu_d \end{pmatrix}$$

- Rotate – Project to the new axes using a dot product $(x_i - \mu) \cdot a_j$

PCA - Principle Component Analysis




- Projection Matrix

- A full projection is defined by a matrix in which each column is a vector defining the direction of one of the new axes

$$A = [a_1, a_2, \dots, a_k] \in \mathbb{R}^{d \times k}$$

- We further require that these vectors will be orthogonal, and often also orthonormal (of a unit length), thus $A^T A = I$




PCA – Finding the first PC

- Let $x'_1 = a_1^T X^T$ be the data projected on the first PC, a_1

$$\begin{aligned} \text{Var}[x'_1] &= E[(x'_1 - \mu'_1)^2] = E[x'^2_1] \\ &= E[(a_1^T X^T)^2] = a_1^T E[X^T X] a_1 = a_1^T S a_1 \end{aligned}$$

- Where S (sometimes notated Σ) is the covariance matrix of X
- We want to find a_1 that maximizes $\text{Var}[x'_1]$ subject to the constraint $a_1^T a_1 = 1$
- This is optimization problem...



PCA – Finding the first PC

- We use Lagrange to find the max of:

$$a_1^T S a_1 - \lambda(a_1^T a_1 - 1)$$

- The derivative with respect to a_1 :

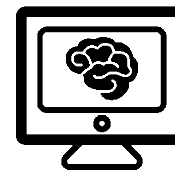
$$2S a_1 - 2\lambda a_1 = 0$$

$$S a_1 = \lambda a_1$$

- This means that a_1 is the eigenvector of S and λ is the corresponding eigenvalue
- Furthermore, we can substitute the result in the original equation:

$$\text{Var}[x'_1] = a_1^T S a_1 = a_1^T \lambda a_1 = \lambda a_1^T a_1 = \lambda$$

- And we got that $\max(\text{Var}[x'_1]) = \max(\lambda)$



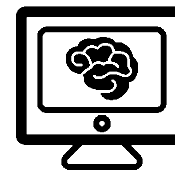
PCA – Finding the PC's

- We can prove that the last equation hold for all PC's, and get:

$$\text{Var}[x'_i] = \lambda_i$$

$$\max(\text{Var}[x'_i]) = \max(\lambda_i)$$

- The PC's are the eigenvectors of S (the covariance matrix of X) and the λ are the corresponding eigenvalues
- The PC's that preserve most of the information are those with the largest λ



PCA – algorithm

- Build the covariance matrix S
- Find Eigenvectors and Eigenvalues of S by solving $(S - \lambda_i I)a_i = 0$
- Sort Eigenvectors by Eigenvalues
- Build the matrix A_j using j eigenvectors with the j greatest eigenvalues
- Transform $x' = A_j(x - \mu)$



PCA – example

- 2 features space, 10 instances

	x_1	x_2
	1.4	1.65
	1.6	1.975
	-1.4	-1.775
	-2	-2.525
	-3	-3.95
	2.4	3.075
	1.5	2.025
	2.3	2.75
	-3.2	-4.05
	-4.1	-4.85
mean	-0.45	-0.5675
var	6.422778	9.952785

Mean
→

	$x_1 - \mu_1$	$x_2 - \mu_2$
	1.85	2.2175
	2.05	2.5425
	-0.95	-1.2075
	-1.55	-1.9575
	-2.55	-3.3825
	2.85	3.6425
	1.95	2.5925
	2.75	3.3175
	-2.75	-3.4825
	-3.65	-4.2825
mean	0	0
var	6.422778	9.952785

- The covariance matrix

$$S = X^T X = \begin{pmatrix} 6.4228 & 7.9876 \\ 7.9876 & 9.9528 \end{pmatrix}$$



PCA – example

- Finding the eigenvalues:

$$\det(S - \lambda_i I) = 7.9876^2 - (6.4228 - \lambda)(9.9528 - \lambda) = 0$$

$$\lambda^2 - 16.3756\lambda + 0.123 = 0$$

- We get $\lambda = 16.3681, 0.0075$
 - Note $\lambda_1 + \lambda_2 = 16.3756 = 6.4228 + 9.9528 = \text{Var}(x_1) + \text{Var}(x_2)$



PCA – example

- Finding the eigenvectors:

$$(S - \lambda_i I) a_i = 0$$
$$\begin{pmatrix} 6.4228 - 16.3681 & 7.9876 \\ 7.9876 & 9.9528 - 16.3681 \end{pmatrix} \begin{pmatrix} a_1^{(1)} \\ a_1^{(2)} \end{pmatrix} = 0$$
$$\begin{pmatrix} -9.9453 & 7.9876 \\ 7.9876 & -6.4153 \end{pmatrix} \begin{pmatrix} a_1^{(1)} \\ a_1^{(2)} \end{pmatrix} = 0$$

- We get

$$a_1^{(1)} = 0.6262, a_1^{(2)} = 0.7797$$

- Note that $a_1^T a_1 = 1$

- We find a_2 with λ_2 , and get

$$a_2^{(1)} = 0.7797, a_2^{(2)} = -0.6262$$



PCA – example

- Build A_j and transform $x' = A_j(x - \mu)$

	$x_1 - \mu_1$	$x_2 - \mu_2$
	1.85	2.2175
	2.05	2.5425
	-0.95	-1.2075
	-1.55	-1.9575
	-2.55	-3.3825
	2.85	3.6425
	1.95	2.5925
	2.75	3.3175
	-2.75	-3.4825
	-3.65	-4.2825
mean	0	0
var	6.422778	9.952785

$$A_j = \begin{pmatrix} 0.6262 & 0.7797 \\ 0.7797 & -0.6262 \end{pmatrix}$$




→

	x'_1	x'_2
	2.887455	0.053846
	3.266097	0.006272
	-1.53638	0.015422
	-2.49687	0.017252
	-4.23415	0.129887
	4.624727	-0.05879
	3.242462	-0.10301
	4.308705	0.066756
	-4.43736	0.036567
	-5.6247	-0.1642
mean	0	0
var	16.3681	0.0075


- So, if we measure the information with VAR, we can see that the first PC has

$$\frac{16.3681}{16.3681 + 0.0075} = 0.999$$

PCA – example visualization



PCA – classification problem



Questions

